


Message Passing



6 Basic Functions

MPI_Init ()

MPI_Comm_rank (Comm)

MPI_Comm_size (Comm)

MPI_Send (Data, Num, Type, Dest, Tag, Comm)

MPI_Recv (Data, Num, Type, Src, Tag, Comm, Stat)

MPI_Finalize ()

Communication type

Point-to-Point communication

Collective communication

Blocking communication

Non-blocking communication

Non-blocking communication

`MPI_Isend (Data, Num, Type, Dest, Tag, Comm, &Req)`

`MPI_Irecv (Data, Num, Type, Src, Tag, Comm, &Req)`

`MPI_Wait (&Req, &Status)`

`MPI_Waitall (Num, Req[], Stat[])`

Collective Communication

MPI_Bcast

MPI_Ibcast

MPI_Scatter

MPI_Gather

MPI_Allscatter

MPI_Alltoall

Collective Communication

rank	send buf		recv buf
----	-----		-----
0	a,b,c	MPI_Allgather ----->	a,b,c,A,B,C,#,@,%
1	A,B,C		a,b,c,A,B,C,#,@,%
2	#, @, %		a,b,c,A,B,C,#,@,%

rank	send buf		recv buf
----	-----		-----
0	a,b,c	MPI_Alltoall ----->	a,A,#
1	A,B,C		b,B,@
2	#, @, %		c,C,%

(a more elaborate case with two elements per process)

rank	send buf		recv buf
----	-----		-----
0	a,b,c,d,e,f	MPI_Alltoall ----->	a,b,A,B,#,@
1	A,B,C,D,E,F		c,d,C,D,%,\$
2	#, @, %, \$, &, *		e,f,E,F,&,*

Example

Rank i wants to send message i to all other ranks

Implementation 1 : Blocking one-to-one:

```
for (int irank = 0; irank < size; ++irank) {  
    if (irank == rank) continue;  
    int stag = irank + size * rank;  
    MPI_Send (&sendbuff, 1, MPI_INT, irank, stag, MPI_COMM_WORLD);  
    std::cout << "rank " << rank << " sending message " << sendbuff << " to rank " << irank << std::endl;  
    int rtag = rank + size * irank;  
    MPI_Recv (&(recvbuff[irank]), 1, MPI_INT, irank, rtag, MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
}
```

Notice encoding of sender and receiver in tag!

DANGER : Deadlock!

Example

Rank i wants to send message i to all other ranks

Implementation 2 : Non-Blocking one-to-one:

```
for (int irank = 0; irank < size; ++irank) {  
    if (irank == rank) continue;  
    int stag = irank + size * rank;  
    MPI_Isend (&sendbuff, 1, MPI_INT, irank, stag, MPI_COMM_WORLD, &(reqs[numreq++]));  
    std::cout << "rank " << rank << " sending message " << sendbuff << " to rank " << irank << std::endl;  
    int rtag = rank + size * irank;  
    MPI_Irecv (&(recvbuff[irank]), 1, MPI_INT, irank, rtag, MPI_COMM_WORLD, &(reqs[numreq++]));  
}  
MPI_Waitall (numreq, &(reqs[0]), MPI_STATUSES_IGNORE);
```

Notice encoding of sender and receiver in tag!

Notice Use of Waitall!

Example

Rank i wants to send message i to all other ranks

Implementation 3 : Blocking Collective (All-to-all)

```
recvbuff.assign (recvbuff.size (), rank);  
MPI_Alltoall (MPI_IN_PLACE, 1, MPI_INT, &(recvbuff[0]), 1, MPI_INT, MPI_COMM_WORLD);
```

Notice MPI_IN_PLACE

Very convenient when messages of same known size are to be sent / received!

Example

Rank i wants to send message i to all other ranks

Implementation 4 : Blocking Collective (All-gather)

```
recvbuff.assign (recvbuff.size (), rank);  
MPI_Allgather (&sendbuff, 1, MPI_INT, &(recvbuff[0]), 1, MPI_INT, MPI_COMM_WORLD);
```

Very convenient when messages of same known size are to be sent / received!

Advanced Example


Distributed assembly of sparse_matrix :

- Each rank **owns** a contiguous range of rows
- Each rank can **assign** elements within its ownership range or **outside** its ownership range
- If multiple ranks assign one matrix element, after assembly the actual value must be the **sum** of local contributions

This is the most common situation encountered in **Finite Element** problems


- Each rank **owns** a range of mesh **elements** and a range of **DOFs**
- For low order continuous formulations **DOFs** correspond to mesh **vertices**
- The same vertex touches multiple **elements**
- The matrix entries related to a given DOF are the sum of contributions from all elements touching the corresponding vertex

Exercise : assemble FEM Laplacian



	J0	J1	J2	J3
I0	2	-1	-1	
I1	-1	2		-1
I2	-1		2	-1
I3		-1	-1	2

Exercise : assemble FEM Laplacian



	J0	J1	J2	J3
I0	2	-1	-1	
I1	-1	2		-1
I2	-1		2	-1
I3		-1	-1	2